

CROISSANT: Centralized Relational Interface for Web-scale SPARQL Endpoints

Takahiro Komamizu, Toshiyuki Amagasa,
Hiroyuki Kitagawa



University of Tsukuba
Japan

Linked Data (LD)

- Open data paradigm
- Linking facts in open data
 - RDF (Resource Description Framework)



- e.g.,

Subject	Predicate	Object
iiwas:iiwas2017	dc:title	"iiWAS2017"
iiwas:iiwas2017	dbo:location	dbr:Salzburg
dbr:Salzburg	dbo:country	dbr:Austria



Search over LD

- Finding facts in LD data
- Standardized method: SPARQL query
 - Graph pattern-based requirement representation
 - Bindings to variables in patterns are results.




Motivation

- LD is **heterogeneous, complicated, and distributed.**
 - Large number of facts
 - Large number of classes (types of entities)
 - Complicated multi-graph structure
 - Distributed and self-managed data servers
- Users are demanded to describe search intents by SPARQL.
 - Graph patterns
 - ➔ require understanding of the structure of data
 - Multiple queries for distributed data
 - ➔ require integration of multiple results

Our Objective and Approach

- Objective
 - Querable interface for popular query languages for distributed RDF data
- Approach: CROISSANT
 - Centralized interface
 - unified access interface for distributed data
 - Relational view-based interface
 - reducing effort for understanding data
 - Query rewriting from SQL to SPARQL
 - Query optimization

CROISSANT: an overview



Users' point-of-view,
SPARQL endpoints are invisible.

CROISSANT: view definition

$\langle name, schema, endpoint_url, SPARQL_query \rangle$

- *name*: view name
 - e.g., movie
- *schema*: relational schema of the view
 - e.g., (movie_id, title, budget)
- *endpoint_url*: location of SPARQL endpoint
 - e.g., <http://dbpedia.org/sparql>
- *SPARQL_query*: SELECT query of data in the view
 - e.g.,

```
SELECT ?movie_id ?title ?budget
WHERE { ?movie_id rdf:type dbo:Film.
         ?movie_id dc:title ?title.
         ?movie_id dbo:budget ?budget. }
```

CROISSANT: query execution

- Naive execution
 1. Execute SPARQL query of corresponding views and store the results into local database.
 2. Perform SQL query over the local database.

SELECT title
FROM movie
WHERE budget > 100,000

Input SQL

title


Results

SELECT ?movie_id ?title ?budget
WHERE { ?movie_id rdf:type dbo:Film.
?movie_id dc:title ?title.
?movie_id dbo:budget ?budget. }

View SPARQL query

movie_id	title	budget

Local database



SPARQL
Endpoint

Performance issues

- Execution cost of SPARQL queries
 - Immature performance of SPARQL endpoints
- Transportation cost from SPARQL endpoints

SELECT title
FROM movie
WHERE budget > 100,000

Input SQL

title


Results

SELECT ?movie_id ?title ?budget
WHERE { ?movie_id rdf:type dbo:Film.
?movie_id dc:title ?title.
?movie_id dbo:budget ?budget. }

View SPARQL query

movie_id	title	budget

Local database



SPARQL
Endpoint

Query Optimization

- Basic idea:
reduce #results from SPARQL endpoints
- Strategies
 - View materialization
 - Everything is transferred to local database in advance.
 - Suffer from update issues.
 - Projection push-down
 - Selection push-down
 - View query merge

Projection Push-down

- Push projection conditions into view SPARQL queries.

```
SELECT title
FROM movie
WHERE budget > 100,000
```

Input SQL

```
SELECT ?movie_id ?title ?budget
WHERE { ?movie_id rdf:type dbo:Film.
         ?movie_id dc:title ?title.
         ?movie_id dbo:budget ?budget. }
```

Pushed-down View SPARQL query

Selection Push-down

- Push selection conditions into view SPARQL queries.
- Pushing rules

Comparator	Variable type	FILTER expression
θ	Numeric	<code>FILTER (variable θ value)</code>
=	Textual	<code>FILTER (str(variable) = value)</code>
<i>like</i>	Textual	<code>FILTER regex(variable, value)</code>

```
SELECT title  
FROM movie  
WHERE budget > 100,000
```


Input SQL

```
SELECT ?movie_id ?title ?budget  
WHERE { ?movie_id rdf:type dbo:Film.  
        ?movie_id dc:title ?title.  
        ?movie_id dbo:budget ?budget.  
FILTER (?budget > 100,000) }
```

Pushed-down View SPARQL query


Naive Join Query Processing

- Steps
 1. Materialize joining views
 2. Perform relational join



View Query Merge – motivation

- Naive join is not efficient if two views are of same SPARQL endpoints.
 - Two separate SPARQL queries are performed.
 - Two sets of larger results are transferred.



View Query Merge – approach

- Idea:
Combine two view SPARQL query on the same SPARQL endpoint into a single query.
- Approach:
 1. Combine projection variables into one set.
 2. Combine graph patterns into one pattern set.
 3. Put FILTER clause for join conditions.
 - If a condition is equality of same attribute names, corresponding FILTER clause is eliminated.

View Query Merge – example

Input SQL

```
SELECT m.title, ma.actor  
FROM movie m, movie_actor ma  
WHERE m.movie_id = ma.movie_id
```

```
SELECT ?movie_id ?title ?budget  
WHERE { ?movie_id rdf:type dbo:Film.  
        ?movie_id dc:title ?title.  
        ?movie_id dbo:budget ?budget. }
```

Movie view

```
SELECT ?movie_id ?actor  
WHERE { ?movie_id rdf:type dbo:Film.  
        ?movie_id dbo:starring ?actor. }
```

Movie-actor view

```
SELECT ?movie_id ?title ?budget ?actor  
WHERE { ?movie_id rdf:type dbo:Film.  
        ?movie_id dc:title ?title.  
        ?movie_id dbo:budget ?budget.  
        ?movie_id rdf:type dbo:Film.  
        ?movie_id dbo:starring ?actor.  
        FILTER (movie_id = movie_id)  }
```

Merged SPARQL query

Experimental Evaluation

- Objective: Check efficiency of CROISSANT
 - Efficiency: query execution time
- Dataset: DBpedia (<http://dbpedia.org/sparql>)
- Views: movie, actor, movie_actor
- Queries (on the same SPARQL endpoint)
 - Two selection queries of different selectivity
 - Observe effect of selectivity
 - Two join queries w/ and w/o selection condition
 - Observe effect of view query merge
 - Observe overall performance of CROISSANT

Experimental Results

Bars

P: Pure execution

M: Materialization


SP: Selection push-down

MG: View query marge

MGSP: SP + MP

Insights

- Materialization is obviously the best.
- Selection push-down works when low selectivity.
- View query merge is **powerful** for join queries.



Conclusion & Future Work

- CROISSANT
 - Centralized interface
 - Relational view-based interface
 - Query rewriting from SQL to SPARQL
 - Query optimization
 - Experimental evaluation introduces optimization techniques work.
- Future work
 - Update issue for taking full advantage of materialization
 - Unified optimization with materialization